eRisk 2026: T1, T2 & T3 results

Anxo Perez[1], Javier Parapar[1], Xi Wang[2], and Fabio Crestani[3]

[1] Information Retrieval Lab, Centro de Investigación en Tecnoloxías da Información e as Comunicacións (CITIC), Universidade da Coruña. Campus de Elviña s/n, C.P. 15071 A Coruña, Spain `javier.parapar@udc.es, anxo.pvila@udc.es`

[2] Department of Computer Science, University of Sheffield, Sheffield, England, United Kingdom `xi.wang@sheffield.ac.uk`

[3] Faculty of Informatics, Università della Svizzera italiana (USI). Campus EST, Via alla Santa 1, 6900 Viganello, Switzerland `fabio.crestani@usi.ch`

Abstract This is the preliminary report of results for preparing the working notes.

1 Task 1: Conversational Depression Detection via LLMs (1st edition)

This task builds on the pilot experience introduced in 2025, where participants interacted with LLM-based personas to detect signs of depression through conversation. In 2026, the task was consolidated as a first official edition, with several changes aimed at improving accessibility, reproducibility, and comparability across participating systems. In particular, the simulated personas were released as LoRA adapters on top of `Meta-Llama-3-8B-Instruct` through Hugging Face, allowing teams to download the models, conduct their interactions locally, and submit their predictions after a limited evaluation period.

The main objective of the task is to determine whether each LLM persona exhibits signs of depression and, if so, to estimate both the **overall depression severity** and the **active depressive symptoms** expressed during the conversation. The personas were fine-tuned using diverse user histories and designed to simulate natural, coherent conversational behaviour. They reflect different depression severity profiles guided by the BDI-II questionnaire, enabling the evaluation of systems across a spectrum of simulated cases, including both depressed and control personas.

Unlike traditional classification tasks based on a fixed set of writings, this task requires participants to actively interact with the persona before making a decision. Teams were allowed to conduct multiple exchanges with each model, but the evaluation setting encouraged early and accurate decisions: systems that required a large number of conversational turns before submitting a prediction could be penalized. Therefore, the task evaluates not only the final diagnostic inference but also the ability to obtain relevant evidence efficiently through conversation.

Participants were not expected to ask the persona directly whether they had depression or whether they were experiencing specific mental health problems.

The personas were explicitly designed to avoid answering such direct questions and could respond with discomfort when confronted with overly explicit inquiries about their mental health. Instead, systems had to infer the possible condition by analysing the persona's language, tone, expressed thoughts, and reactions throughout the interaction. This setup encourages the development of conversational strategies that are sensitive, indirect, and clinically informative.

Each team could submit up to three runs using different methods. At most one run could involve manual intervention or human assistance, while the remaining runs had to be fully automated. Moreover, each submitted method was required to operate independently, without sharing information across runs. The evaluation focused on two complementary aspects: (i) the accurate identification of depressive symptoms expressed by the persona, and (ii) the estimation of the overall depression level according to BDI-II standards.

The BDI-II score ranges from 0 to 63, with higher scores indicating greater depression severity. These scores are commonly grouped into four severity levels: *minimal depression* (0–9), *mild depression* (10–18), *moderate depression* (19–29), and *severe depression* (30–63). In addition to estimating the overall severity level, participants were asked to identify which of the 21 BDI-II symptoms were present in the persona's behaviour, as listed in Table 1.

Table 1. The 21 BDI-II Depression Symptoms.

21 Depression Symptoms		
Sadness	Pessimism	Past Failure
Loss of Pleasure	Guilty Feelings	Punishment Feelings
Self-Dislike	Self-Criticalness	Suicidal Thoughts or Wishes
Crying	Agitation	Loss of Interest
Indecisiveness	Worthlessness	Loss of Energy
Changes in Sleeping Pattern	Irritability	Changes in Appetite
Concentration Difficulty	Tiredness or Fatigue	Loss of Interest in Sex

1.1 Task 1: Participant teams

Table 2 shows the participating teams and several statistics about their interactions with the LLM personas, including the number of submitted runs, the number of personas covered in each run, whether any run involved manual intervention, the mean number of messages per conversation, and the mean number of characters per message. Overall, the numbers reveal a diverse set of interaction strategies, ranging from short and direct conversations to longer multi-turn dialogues. In this table, messages correspond only to those sent by the participant's system to the LLM persona.

A total of 21 teams participated in this first edition of the task. Most teams submitted multiple runs, with 16 teams submitting three runs. Several teams covered the complete set of 20 personas in all their submitted runs, namely

Table 2. Task 1: participating teams, number of runs, number of personas submitted per run, presence of manual intervention, mean number of messages per conversation, and mean number of characters per message.

Team	#Runs	Personas/run	Manual run(s)	Interaction statistics msgs./conv.	chars./msg.
Amor-Fati	3	20; 16; 13	#1	9.3	108.76
Awakened	2	20; 20	–	5.3	84.37
BUAP-CRC	3	20; 20; 20	#3	13.2	140.82
DS-GT	3	20; 20; 20	–	11.5	210.15
INSA-Lyon	3	19; 19; 19	–	6.2	156.96
INSALyon-2	3	19; 19; 19	–	23.4	80.41
NoirByte	1	12	–	4.6	97.71
PMH-PsychAI-Lab	1	19	..	8.0	131.16
RecEarlyMind	1	20	–	7.5	93.23
SINAI	3	20; 20; 20	–	4.4	242.61
SVNITDS100	3	6; 6; 6	#1	4.9	114.51
SymptomAI	3	2; 2; 8	–	–	119.41
Techbbas	3	6; 6; 6	–	8.0	186.49
The-Interrogators	3	6; 6; 6	#1	9.3	206.94
UNED-GELP	3	20; 20; 18	–	10.7	114.47
VANGUARD	3	2; 2; 18	–	15.0	80.62
Xtra	1	2	–	6.0	192.46
erisk-cedri	3	18; 18; 18	–	4.8	82.79
lexxy	3	20; 20; 20	–	29.4	93.84
pjmathematician	3	20; 20; 20	–	6.0	108.88
upb-uottawa	3	20; 20; 20	#1	25.6	119.85

Awakened, BUAP-CRC, DS-GT, SINAI, lexxy, pjmathematician, and upb-uottawa. Other teams submitted almost complete runs, such as INSA-Lyon and INSALyon-2, with 19 personas per run, or erisk-cedri, with 18 personas per run. By contrast, teams such as SVNITDS100, SymptomAI, Techbbas, The-Interrogators, VANGUARD, and Xtra submitted partial runs, suggesting either exploratory participation, incomplete processing of the released personas, or selective evaluation strategies.

The interaction statistics show substantial variation in the number of messages exchanged before producing a decision. The most extensive conversational strategies were followed by lexxy and upb-uottawa, with averages of 29.4 and 25.6 participant messages per conversation, respectively. INSALyon-2 also relied on relatively long dialogues, with 23.4 messages on average, while VANGUARD used 15.0 messages per conversation. At the other end of the spectrum, SINAI, NoirByte, SVNITDS100, erisk-cedri, and Awakened used shorter interactions, all below five and a half participant messages per conversation on average. These differences suggest that participants explored substantially different trade-offs between early decision-making and evidence gathering.

The mean number of characters per message also varied considerably. SINAI produced the longest messages on average, with 242.61 characters per message, followed by DS-GT, The-Interrogators, Xtra, and Techbbas, all above 180 characters per message. These runs appear to favour more detailed prompts or information-rich conversational turns. Conversely, teams such as INSALyon-2, VANGUARD, erisk-cedri, and Awakened used shorter messages, below 85 characters per turn on average, suggesting more concise or template-based interaction strategies. For SymptomAI, the mean number of messages per conversation is not reported, although the average message length is included.
Finally, several teams included manual or human-assisted components in one of their runs, as allowed by the task rules. This was the case for Amor-Fati, BUAP-CRC, SVNITDS100, The-Interrogators, and upb-uottawa. The inclusion of both automated and manual-in-the-loop submissions makes it possible to compare fully automatic conversational strategies against approaches where human judgement guided part of the interaction or decision process.

1.2 Evaluation Metrics

Based on evaluation metrics that have been developed from eRisk 2019 [2], which involved the use of BDI-II questionnaires and scores, we extend and develop the evaluation approaches as follows:

- **Depression Category Hit Rate (DCHR)**: Based on the four depression level categories that we have discussed, from minimal depression to severe depression, this effectiveness measure examines the fraction of cases where the BDI-II scores describing simulated personas estimated by the participants lie in the correct depression category.
- **Average DODL (ADODL)**: For this pilot task, we reuse the Average Difference between Overall Depression Levels (ADODL), which measures the closeness between the actual and estimated depression level for effectiveness measurement. The ADODL is calculated by following: $CR = (MAD - |ADL - EDL|)/MAD$, where $|ADL - EDL|$ calculates the absolute value between the Actual Depression Level (ADL) and the Estimated Depression Level (EDL). Then divided by Maximum Absolute Difference (i.e., 63) to obtain a normalised evaluation score in [0,1]. For example, if a simulated persona has a minor depression severity (depression level 5) and a participant estimates the depression level is 9, the DODL is calculated as $(63 - |9 - 5|)/63 = 0.9365$.
- **Average Symptom Hit Rate (ASHR)**: For the last effectiveness measure, aside from estimating the depression level of simulated personas as per BDI-II scores, this pilot task also involves the identification of major depression symptoms of simulated personas. Hence, SHR calculates the ratio of cases where the participants can correctly identify the major symptoms of the simulated personas. For example, each simulated persona has four major symptoms. If a participant accurately identifies two of them, then the SHR equals $2/4 = 0.5$.

Latency-aware Evaluation In addition to the standard effectiveness metrics, we introduce two latency-aware variants that penalize late decisions. In this task, participants could interact with each LLM persona for an arbitrary number of messages before submitting their final prediction. Therefore, systems that reach accurate conclusions after fewer user messages should be rewarded over systems that require long conversations to obtain the same result.
For each persona $u \in U$, let k_u be the number of messages sent by the participant's system to the LLM persona before producing the final prediction. Following the latency penalty used in early risk detection, we define:

$$penalty(k_u) = -1 + \frac{2}{1 + \exp^{-p \cdot (k_u - 1)}}, \tag{1}$$

where p controls how quickly the penalty increases with the number of messages. A decision made after the first user message receives no penalty, i.e., $penalty(1) = 0$. The corresponding speed factor is:

$$speed(k_u) = 1 - penalty(k_u). \tag{2}$$

Thus, $speed(k_u)$ is close to 1 for early decisions and decreases as the number of messages increases.
In this first edition, we set p so that the speed factor is equal to 0.5 after $K_0 = 30$ user messages. This value was chosen based on the observed interaction statistics of the submitted runs: many systems used approximately 10–30 messages per conversation, while a smaller number of systems relied on substantially longer interactions. Therefore, $K_0 = 30$ provides a balanced threshold that rewards early decisions while avoiding an excessively harsh penalty for systems that conduct moderately long exploratory conversations.
More formally, we set:

$$p = \frac{\ln(3)}{K_0 - 1} = \frac{\ln(3)}{30 - 1} \simeq 0.0379. \tag{3}$$

With this setting, a decision made after the first user message receives the full effectiveness score ($speed(1) = 1$), while a decision made after 30 user messages receives half of the original score ($speed(30) = 0.5$). Decisions made after longer conversations receive progressively lower speed factors.
Let $DCHR_u$ be the persona-level Depression Category Hit Rate score, where $DCHR_u = 1$ if the estimated BDI-II score falls into the correct depression category and $DCHR_u = 0$ otherwise. We define the latency-weighted category hit rate as:

$$LDCHR = \frac{1}{|U|} \sum_{u \in U} DCHR_u \cdot speed(k_u). \tag{4}$$

Similarly, let SHR_u be the persona-level Symptom Hit Rate, computed as the proportion of key symptoms correctly identified for persona u. We define the latency-weighted average symptom hit rate as:

$$LASHR = \frac{1}{|U|} \sum_{u \in U} SHR_u \cdot speed(k_u). \tag{5}$$

These two metrics preserve the interpretation of the original measures while incorporating the cost of delayed interaction. A system that correctly identifies the depression category or symptoms after a short conversation will obtain latency-aware scores close to its original DCHR or ASHR. Conversely, a system that reaches the same prediction only after many messages will receive lower latency-aware scores. In this way, $LDCHR$ and $LASHR$ evaluate not only the correctness of the final prediction, but also the efficiency of the conversational strategy used to obtain it.

1.3 Task 1: Results

The results are reported separately for complete and incomplete submissions. Table 3 includes only those runs that submitted predictions for all 20 LLM personas, while Table 4 reports runs that covered fewer than 20 personas. This distinction is important because the metrics are computed over the personas submitted by each run. Therefore, scores obtained from a very small subset of personas are not directly comparable to those obtained over the full evaluation set. The complete-submission table provides the most reliable basis for comparing systems under the full task setting, whereas the incomplete-submission table is useful for documenting partial participation and exploratory runs. In addition to the original effectiveness metrics, both tables include the latency-aware variants LDCHR and LASHR, which penalize correct decisions that are reached only after longer conversations. Each metric value is also accompanied by its rank within the corresponding table, allowing the reader to identify not only the best overall run but also the relative position of each run for each individual metric.

For complete submissions, the best ADODL score was achieved by PJMATHEMATICIAN in run #3, with an ADODL of 0.9254. The strongest DCHR result was obtained by BUAP-CRC in run #2, with a score of 0.6000, while the best ASHR was achieved by UPB-UOTTAWA in run #3, with 0.3375. When considering the more restrictive latency-aware metrics, with $K_0 = 10$, PJMATHEMATICIAN remained the strongest team: runs #1 and #3 obtained the highest LDCHR value, 0.3872, and run #2 achieved the highest LASHR value, 0.2288. This indicates that PJMATHEMATICIAN combined strong prediction quality with comparatively efficient conversational strategies. However, the latency-aware rankings also show that some systems with good original scores were penalized substantially when they required longer interactions. For example, LEXXY and UPB-UOTTAWA achieved competitive ASHR values, but their LASHR ranks dropped due to longer conversations. Overall, pjmathematician, DS-GT, and BUAP-CRC obtained the highest ADODL values among full-coverage submissions, while the LDCHR and LASHR columns highlight the additional importance of reaching accurate decisions early.

For incomplete submissions, some runs obtained very high values, such as SymptomAI run #2 with an ADODL of 0.9683 and VANGUARD run #2 with a DCHR of 1.0000. However, both results were computed over only 2 out of the 20 personas, and should therefore be interpreted with caution. The same caveat applies to the latency-aware metrics: although Xtra run #1 obtained the highest LASHR value, 0.1760, and VANGUARD run #2 obtained a strong LDCHR value, both runs covered only two personas. Among the partial submissions with broader coverage, UNED-GELP run #2 covered 18 personas and achieved an ADODL of 0.8836 and the best LDCHR value in the incomplete table, 0.4550. Similarly, erisk-cedri submitted three 18-persona runs, with its run #3 obtaining an ADODL of 0.8686 and an LDCHR of 0.2817. The two INSA-Lyon variants covered 19 personas per run, with ADODL values mostly in the 0.83–0.85 range, although their latency-aware scores varied depending on the length of the conversations. These results show that several incomplete submissions were close to the full-coverage setting, whereas others correspond to very small subsets and mainly serve as partial evidence of system behaviour.

Table 3. Evaluation for Task 1 considering only runs that submitted predictions for all 20 LLM personas. Teams are ordered by their best ADODL run. Ranks for each metric are shown in parentheses. '*' indicates manual or human-assisted runs.

Team	Run	DCHR	ADODL	ASHR	LDCHR	LASHR
pjmathematician	#1	0.5500 (2)	0.9071 (2)	0.3125 (3)	**0.3872 (1)**	0.2200 (2)
	#2	0.4500 (5)	0.9040 (5)	0.3250 (2)	0.3168 (3)	**0.2288 (1)**
	#3	0.5500 (2)	**0.9254 (1)**	0.3000 (4)	**0.3872 (1)**	0.2112 (3)
DS-GT	#1	0.4500 (5)	0.8841 (11)	0.2500 (8)	0.2091 (10)	0.1138 (8)
	#2	0.4000 (6)	0.8968 (8)	0.2500 (8)	0.1696 (15)	0.1142 (7)
	#3	0.5500 (2)	0.9063 (3)	0.1875 (11)	0.2493 (6)	0.0777 (14)
BUAP-CRC	#1	0.5000 (3)	0.9056 (4)	0.1000 (16)	0.2028 (11)	0.0310 (19)
	#2	**0.6000 (1)**	0.9024 (6)	0.1500 (15)	0.2409 (7)	0.0575 (15)
	#3*	0.4500 (5)	0.8929 (9)	0.1750 (12)	0.1761 (14)	0.0548 (16)
SINAI	#1	0.5000 (3)	0.8905 (10)	0.1875 (11)	0.2500 (5)	0.0938 (11)
	#2	0.4000 (6)	0.8738 (13)	0.2250 (10)	0.2000 (13)	0.1125 (9)
	#3	0.4737 (4)	0.8989 (7)	0.1711 (13)	0.2368 (8)	0.0855 (12)
UNED-GELP	#0	0.4000 (6)	0.8040 (22)	0.1875 (11)	0.0572 (16)	0.0268 (20)
	#1	0.4500 (5)	0.8968 (8)	0.1625 (14)	0.3168 (3)	0.1144 (6)
Awakened	#1	0.2500 (9)	0.8325 (17)	0.1000 (16)	0.2167 (9)	0.0807 (13)
	#2	0.5000 (3)	0.8802 (12)	0.2625 (7)	0.3374 (2)	0.1709 (4)
lexxy	#1	0.3000 (8)	0.8556 (14)	0.1875 (11)	0.0137 (22)	0.0122 (24)
	#2	0.3000 (8)	0.8492 (15)	0.2500 (8)	0.0144 (21)	0.0162 (23)
	#3	0.3000 (8)	0.8413 (16)	0.2875 (5)	0.0186 (20)	0.0196 (22)
Amor-Fati	#1*	0.4500 (5)	0.8254 (18)	0.2375 (9)	0.2514 (4)	0.1410 (5)
upb-uottawa	#1*	0.3500 (7)	0.8222 (20)	0.2750 (6)	0.0376 (17)	0.0315 (18)
	#2	0.3000 (8)	0.8230 (19)	0.2625 (7)	0.0304 (18)	0.0355 (17)
	#3	0.3000 (8)	0.7849 (23)	**0.3375 (1)**	0.0214 (19)	0.0241 (21)
RecEarlyMind	#1	0.3000 (8)	0.8071 (21)	0.1625 (14)	0.2013 (12)	0.0977 (10)

Table 4. Evaluation for Task 1 considering runs that submitted predictions for fewer than 20 LLM personas. Teams are ordered by their best ADODL run. Ranks for each metric are shown in parentheses. '*' indicates manual or human-assisted runs.

Team	Run	Personas	DCHR	ADODL	ASHR	LDCHR	LASHR
SymptomAI	#1	2/20	0.0000 (18)	0.8810 (6)	0.1250 (12)	0.0000 (25)	0.0625 (18)
	#2	2/20	0.5000 (5)	**0.9683 (1)**	0.2500 (3)	0.2500 (7)	0.1250 (6)
	#3	8/20	0.3750 (8)	0.8294 (22)	0.2500 (3)	0.1875 (13)	0.1250 (6)
VANGUARD	#1	2/20	0.5000 (5)	0.8810 (6)	0.1250 (12)	0.1533 (17)	0.0383 (23)
	#2	2/20	**1.0000 (1)**	0.9286 (2)	0.0000 (18)	0.3066 (3)	0.0000 (26)
	#3	18/20	0.2222 (13)	0.7319 (28)	0.1806 (6)	0.0681 (21)	0.0554 (19)
SVNITDS100	#1	6/20	0.2857 (11)	0.8594 (10)	0.1429 (11)	0.1969 (12)	0.0714 (17)
	#2	6/20	0.5000 (5)	0.9259 (3)	0.0417 (16)	0.2500 (7)	0.0208 (25)
	#3*	5/20	0.2000 (15)	0.9111 (4)	0.1500 (9)	0.1521 (18)	0.1199 (8)
UNED-GELP	#2	18/20	0.5556 (4)	0.8836 (5)	0.1667 (7)	**0.4550 (1)**	0.1365 (5)
Techbbas	#1	6/20	0.3333 (9)	0.8280 (23)	0.1667 (7)	0.1990 (11)	0.0995 (12)
	#2	6/20	0.1667 (16)	0.7646 (26)	0.0000 (18)	0.0995 (19)	0.0000 (26)
	#4	6/20	0.1667 (16)	0.8730 (7)	0.1250 (12)	0.0995 (19)	0.0746 (16)
The-Interrogators	#1*	6/20	0.6000 (3)	0.8667 (9)	0.0500 (15)	0.3000 (4)	0.0250 (24)
	#2	6/20	0.6000 (3)	0.8413 (18)	0.0500 (15)	0.3000 (4)	0.0250 (24)
	#3	6/20	0.8000 (2)	0.8730 (7)	0.0500 (15)	0.4000 (2)	0.0250 (24)
erisk-cedri	#1	18/20	0.1111 (17)	0.8325 (20)	0.1667 (7)	0.0878 (20)	0.1200 (7)
	#2	18/20	0.2222 (13)	0.8474 (14)	0.1806 (6)	0.1724 (16)	0.1496 (3)
	#3	18/20	0.3889 (6)	0.8686 (8)	0.0972 (14)	0.2817 (5)	0.0758 (14)
NoirByte	#1	12/20	0.3333 (9)	0.8585 (11)	0.1042 (13)	0.2730 (6)	0.0841 (13)
Amor-Fati	#2	16/20	0.3750 (8)	0.8552 (12)	0.2656 (2)	0.2219 (9)	0.1385 (4)
	#3	13/20	0.3846 (7)	0.8462 (17)	0.1923 (5)	0.2040 (10)	0.1096 (10)
INSALyon-2	#1	19/20	0.2105 (14)	0.8145 (24)	**0.2763 (1)**	0.0203 (24)	0.0542 (20)
	#2	19/20	0.3158 (10)	0.8463 (16)	0.2500 (3)	0.0679 (22)	0.0513 (21)
	#3	19/20	0.3158 (10)	0.8496 (13)	0.2237 (4)	0.0486 (23)	0.0495 (22)
INSA-Lyon	#1	19/20	0.3158 (10)	0.8388 (19)	0.1579 (8)	0.2285 (8)	0.1142 (9)
	#2	19/20	0.2632 (12)	0.8304 (21)	0.1447 (10)	0.1767 (15)	0.1013 (11)
	#3	19/20	0.2632 (12)	0.8471 (15)	0.2237 (4)	0.1826 (14)	0.1525 (2)
Xtra	#1	2/20	0.0000 (18)	0.8016 (25)	0.2500 (3)	0.0000 (25)	**0.1760 (1)**
PMH-PsychAI-Lab	#1	19/20	0.2105 (14)	0.7327 (27)	0.0091 (17)	0.1724 (16)	0.0754 (15)

2 Task 2: Contextualized Early Detection of Depression (2nd edition)

In 2026, we have the second edition of this task, which introduces a different scenario in depression detection by incorporating full conversational contexts. Whereas earlier eRisk editions released isolated posts authored by a single user, the contextualized early detection task provided the *entire Reddit discussion thread* in which the target user intervened. Consequently, in the training and test phase, systems had access not only to the messages produced by the target user but also to every other contribution in the thread and to the interaction structure that links the messages.
This design was motivated by the observation that the clinical relevance of a message often becomes apparent only when interpreted alongside the surrounding conversation (e.g., a user's seemingly neutral reply may reveal hopelessness when it is a response to criticism or a plea for help). Thus, the task is designed to simulate real-world scenarios where depression detection may rely on analyzing exchanges between multiple participants. This setup presents unique challenges, as systems must consider not only the textual content of individual posts but also the interplay between participants and how this context influences the detection of depressive symptoms. The task was divided into two phases:

- During the training phase, participants worked with the test dataset of this task in 2025, consisting of contextualized user writings from depressed and control users.
- The test phase was carried out interactively. For each target user, the server released a sequence of discussion threads in real time. Each thread constituted a *submission round*. After processing the thread for each target user (including every post written by all interlocutors up to that moment), teams had to return a depression label (positive/negative) and a confidence score before the next thread became available.

2.1 Task 2: Evaluation Metrics

Decision-based Evaluation This evaluation approach uses the binary decisions made by the participating systems for each user. In addition to standard classification measures such as Precision, Recall, and F1 score (computed with respect to the positive class), we also calculate ERDE (Early Risk Detection Error), used in previous editions of the lab. A detailed description of ERDE was presented by Losada and Crestani in [1]. ERDE is an error measure that incorporates a penalty for delayed correct alerts (true positives). The penalty increases with the delay in issuing the alert, measured by the number of user posts processed before making the alert.
Since 2019, we complemented the evaluation report with additional decision-based metrics that try to capture additional aspects of the problem. These metrics try to overcome some limitations of $ERDE$, namely:

- the penalty associated to true positives goes quickly to 1. This is due to the functional form of the cost function (sigmoid).
- a perfect system, which detects the true positive case right after the first round of messages (first chunk), does not get error equal to 0.
- with a method based on releasing data in a chunk-based way (as it was done in 2017 and 2018) the contribution of each user to the performance evaluation has a large variance (different for users with few writings per chunk vs users with many writings per chunk).
- $ERDE$ is not interpretable.

Some research teams have analysed these issues and proposed alternative ways for evaluation. Trotzek and colleagues [4] proposed $ERDE_o^{\%}$. This is a variant of ERDE that does not depend on the number of user writings seen before the alert but, instead, it depends on the *percentage* of user writings seen before the alert. In this way, user's contributions to the evaluation are normalized (currently, all users weight the same). However, there is an important limitation of $ERDE_o^{\%}$. In real life applications, the overall number of user writings is not known in advance. Social Media users post contents online and screening tools have to make predictions with the evidence seen. In practice, you do not know when (and if) a user's thread of messages is exhausted. Thus, the performance metric should not depend on knowledge about the total number of user writings.
Another proposal of an alternative evaluation metric for early risk prediction was done by Sadeque and colleagues [3]. They proposed $F_{latency}$, which fits better with our purposes. This measure is described next.
Imagine a user $u \in U$ and an early risk detection system that iteratively analyzes u's writings (e.g. in chronological order, as they appear in Social Media) and, after analyzing k_u user writings ($k_u \geq 1$), takes a binary decision $d_u \in \{0, 1\}$, which represents the decision of the system about the user being a risk case. By $g_u \in \{0, 1\}$, we refer to the user's golden truth label. A key component of an early risk evaluation should be the delay on detecting true positives (we do not want systems to detect these cases too late). Therefore, a first and intuitive measure of delay can be defined as follows[4]:

$$\text{latency}_{TP} = \text{median}\{k_u : u \in U, d_u = g_u = 1\} \tag{6}$$

This measure of latency is calculated over the true positives detected by the system and assesses the system's delay based on the median number of writings that the system had to process to detect such positive cases. This measure can be included in the experimental report together with standard measures such as Precision (P), Recall (R) and the F-measure (F):

[4] Observe that Sadeque et al (see [3], pg 497) computed the latency for all users such that $g_u = 1$. We argue that latency should be computed only for the true positives. The false negatives ($g_u = 1$, $d_u = 0$) are not detected by the system and, therefore, they would not generate an alert.

$$P = \frac{|u \in U : d_u = g_u = 1|}{|u \in U : d_u = 1|} \tag{7}$$

$$R = \frac{|u \in U : d_u = g_u = 1|}{|u \in U : g_u = 1|} \tag{8}$$

$$F = \frac{2 \cdot P \cdot R}{P + R} \tag{9}$$

Furthermore, Sadeque et al. proposed a measure, $F_{latency}$, which combines the effectiveness of the decision (estimated with the F measure) and the delay[5] in the decision. This is calculated by multiplying F by a penalty factor based on the median delay. More specifically, each individual (true positive) decision, taken after reading k_u writings, is assigned the following penalty:

$$penalty(k_u) = -1 + \frac{2}{1 + \exp^{-p \cdot (k_u - 1)}} \tag{10}$$

where p is a parameter that determines how quickly the penalty should increase. In [3], p was set such that the penalty equals 0.5 at the median number of posts of a user[6]. Observe that a decision right after the first writing has no penalty (i.e. $penalty(1) = 0$). Figure 1 plots how the latency penalty increases with the number of observed writings.

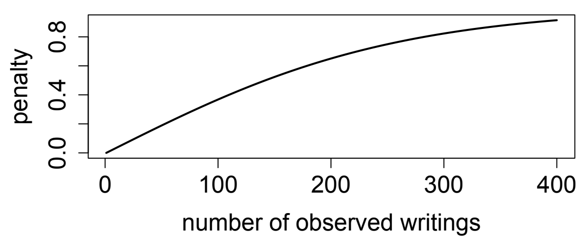

Figure 1. Latency penalty increases with the number of observed writings (k_u)

The system's overall speed factor is computed as:

$$speed = (1 - \text{median}\{penalty(k_u) : u \in U, d_u = g_u = 1\}) \tag{11}$$

[5] Again, we adopt Sadeque et al.'s proposal but we estimate latency only over the true positives.

[6] In the evaluation we set p to 0.0078, a setting obtained from the eRisk 2017 collection.

where speed equals 1 for a system whose true positives are detected right at the first writing. A slow system, which detects true positives after hundreds of writings, will be assigned a speed score near 0.
Finally, the *latency-weighted* F score is simply:

$$F_{latency} = F \cdot speed \tag{12}$$

Since 2019 user's data were processed by the participants in a post by post basis (i.e. we avoided a chunk-based release of data). Under these conditions, the evaluation approach has the following properties:

- smooth grow of penalties;
- a perfect system gets $F_{latency} = 1$;
- for each user u the system can opt to stop at any point k_u and, therefore, now we do not have the effect of an imbalanced importance of users;
- $F_{latency}$ is more interpretable than $ERDE$.

Ranking-based Evaluation In addition to the evaluation discussed above, we employed an alternative form of evaluation to further assess the systems. After each data release (new user writing, that is post or comment), participants were required to provide the following information for each user in the collection:

- A decision for the user (alert or no alert), which was used to calculate the decision-based metrics discussed previously.
- A score representing the user's level of risk, estimated based on the evidence observed thus far.

The scores were used to create a ranking of users in descending order of estimated risk. For each participating system, a ranking was generated at each data release point, simulating a continuous re-ranking approach based on the observed evidence. In a real-life scenario, this ranking would be presented to an expert user who could make decisions based on the rankings (e.g., by inspecting the top of the rankings). Each ranking can be evaluated using standard ranking metrics such as P@10 or NDCG. Therefore, we report the performance of the systems based on the rankings after observing different numbers of writings.

2.2 Task 2: Participant teams

Table 5 reports the teams that participated in Task 2, together with the number of submitted runs, the number of user threads processed, and the lapse of time between each team's first and last response. As in the previous edition, this lapse of time provides an approximate indication of the level of automation and computational efficiency of the submitted systems.
A total of 17 teams participated in this second edition of the contextualized early detection task. Most teams submitted the maximum number of runs allowed,

with 12 teams submitting five runs. In terms of coverage, 10 teams processed the complete set of 500 user threads: erisk-cedri, HUTECH-NLP, Lotu-ixa, UET-PsyWar, DeepCare, UNED-GELP, INSA-Lyon, HUGETIME, LCDAA, and Snugi-AI-v2. The remaining teams processed a partial subset of the test collection, ranging from 4 to 302 user threads.
The fastest complete submissions were produced by Snugi-AI-v2 and erisk-cedri, both completing the task in under two hours, followed by Lotu-ixa, UET-PsyWar, and DeepCare, which completed the full set in under ten hours. These results suggest that, despite the additional complexity introduced by full conversational contexts, efficient end-to-end processing remains feasible. At the other end of the spectrum, teams such as INSA-Lyon, LCDAA, NYCU-NLP, and especially NoirByte required substantially longer processing times, which may reflect more computationally demanding pipelines, manual intervention, or interruptions during the interactive evaluation process. Finally, TeamZed submitted five runs but processed only four user threads, and should therefore be interpreted as a very limited partial submission.

Table 5. Task 2: participating teams, number of runs, number of threads processed by the team, and lapse of time taken for the entire process.

Team	#Runs	#User threads	Lapse of time (from 1st to last response)
Awakened	1	118	0 days 23:12
erisk-cedri	5	500	0 days 01:47
VANGUARD	2	129	0 days 12:42
HUTECH-NLP	5	500	2 days 03:02
Lotu-ixa	5	500	0 days 04:24
UET-PsyWar	5	500	0 days 06:57
DeepCare	5	500	0 days 09:14
UNED-GELP	5	500	0 days 10:27
MacEarlyDetect	5	121	0 days 21:51
INSA-Lyon	5	500	6 days 01:45
NoirByte	5	289	53 days 18:49
NYCU-NLP	5	247	11 days 08:35
HUGETIME	1	500	2 days 20:14
LCDAA	5	500	3 days 12:07
MindAILab	1	302	0 days 01:15
Snugi-AI-v2	5	500	0 days 01:26
TeamZed	5	4	0 days 00:53

2.3 Task 2: Results

Table 6 show the decision-based results of task 2, ordered in terms of best F1. Table 7 shows the ranking-based results, ordered by $NDCG@100$.

Table 6. Task 2 decision-based evaluation ordered in terms of best F_1.

Team	Run	P	R	F_1	$ERDE_5$	$ERDE_{50}$	latencyTP	$speed$	$F_{latency}$
HUGETIME	0	0.80	0.87	**0.83**	0.15	**0.05**	9.00	0.97	**0.81**
UNED-GELP	0	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00
	1	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00
	2	0.79	0.85	0.82	0.11	0.06	4.00	0.99	**0.81**
	3	0.78	0.71	0.75	0.10	0.08	2.00	**1.00**	0.74
	4	0.81	0.69	0.75	0.11	0.08	3.00	0.99	0.74
NoirByte	0	0.88	0.76	0.82	0.15	0.09	20.00	0.93	0.76
	1	0.93	0.70	0.80	0.16	0.10	24.00	0.91	0.73
	2	0.93	0.62	0.74	0.16	0.11	28.50	0.89	0.66
	3	0.94	0.54	0.69	0.17	0.12	27.00	0.90	0.62
	4	0.93	0.42	0.58	0.17	0.13	35.50	0.87	0.50
MindAILab	0	0.84	0.80	0.82	0.16	0.08	16.00	0.94	0.77
INSA-Lyon	0	0.76	0.85	0.80	0.16	0.07	13.00	0.95	0.76
	1	0.69	0.89	0.78	0.15	**0.05**	9.00	0.97	0.75
	2	0.46	0.95	0.62	0.14	0.08	7.50	0.97	0.60
	3	0.63	0.92	0.75	0.16	0.06	11.50	0.96	0.72
	4	0.75	0.85	0.80	0.16	0.07	14.00	0.95	0.76
UET-PsyWar	0	0.62	0.87	0.72	0.15	0.08	15.00	0.95	0.68
	1	0.57	0.88	0.69	0.16	0.08	15.00	0.95	0.65
	2	0.73	0.82	0.77	0.15	0.07	15.00	0.95	0.73
	3	0.76	0.80	0.78	0.14	0.08	14.00	0.95	0.74
	4	0.60	0.88	0.71	0.15	0.07	15.00	0.95	0.68
LCDAA	0	0.66	0.81	0.73	0.11	0.07	3.00	0.99	0.72
	1	0.69	0.81	0.74	0.11	0.07	4.00	0.99	0.74
	2	0.78	0.78	0.78	0.13	0.08	10.00	0.96	0.75
	3	0.70	0.70	0.70	0.09	0.07	**1.00**	**1.00**	0.70
	4	0.72	0.74	0.73	0.11	0.07	2.00	**1.00**	0.73
Lotu-ixa	0	0.37	0.96	0.54	0.13	0.07	3.00	0.99	0.53
	1	0.31	0.97	0.47	0.14	0.08	3.00	0.99	0.46
	2	0.38	0.93	0.54	0.14	0.08	6.00	0.98	0.53
	3	0.63	0.86	0.73	0.14	0.07	12.00	0.96	0.70
	4	0.75	0.77	0.76	0.12	0.07	6.50	0.98	0.74
erisk-cedri	0	0.73	0.77	0.75	0.10	0.07	1.50	**1.00**	0.75
	1	0.78	0.51	0.61	0.13	0.10	4.50	0.99	0.60
	2	0.85	0.44	0.58	0.14	0.11	5.50	0.98	0.57
	3	0.84	0.34	0.48	0.16	0.12	9.00	0.97	0.47
	4	**1.00**	0.11	0.20	0.17	0.15	9.00	0.97	0.19
VANGUARD	0	0.62	0.90	0.74	0.14	**0.05**	6.50	0.98	0.72
	1	0.95	0.45	0.61	0.17	0.11	16.00	0.94	0.58
HUTECH-NLP	0	0.31	0.99	0.48	0.09	0.07	**1.00**	**1.00**	0.48
	1	0.48	0.96	0.64	0.09	**0.05**	3.00	0.99	0.63
	2	0.77	0.70	0.74	0.11	0.06	4.00	0.99	0.73
	3	0.40	0.93	0.56	**0.08**	0.06	**1.00**	**1.00**	0.56
	4	0.43	0.92	0.59	0.09	**0.05**	3.00	0.99	0.58
Snugi-AI-v2	0	0.72	0.71	0.72	0.18	0.08	8.00	0.97	0.70
	1	0.77	0.69	0.73	0.17	0.09	14.00	0.95	0.69
	2	0.71	0.66	0.69	0.18	0.07	8.00	0.97	0.67
	3	0.80	0.65	0.72	0.17	0.07	8.00	0.97	0.70
	4	0.66	0.65	0.65	0.18	0.08	8.00	0.97	0.63
DeepCare	0	0.92	0.51	0.65	0.11	0.09	**1.00**	**1.00**	0.65
	1	0.54	0.89	0.67	0.15	0.06	7.00	0.98	0.66
	2	0.96	0.25	0.40	0.16	0.13	8.00	0.97	0.39
	3	0.92	0.53	0.67	0.16	0.10	10.00	0.96	0.65
	4	0.95	0.57	0.71	0.13	0.09	5.00	0.98	0.70
NYCU-NLP	0	0.43	0.79	0.55	0.16	0.09	11.00	0.96	0.53
	1	0.47	0.60	0.53	0.17	0.11	15.00	0.95	0.50
	2	0.49	0.44	0.46	0.18	0.12	13.50	0.95	0.44
	3	0.49	0.64	0.56	0.17	0.11	15.00	0.95	0.52
	4	0.49	0.64	0.56	0.17	0.11	15.00	0.95	0.52
Awakened	0	0.33	0.68	0.45	0.11	0.10	**1.00**	**1.00**	0.45
MacEarlyDetect	0	0.21	0.98	0.35	0.16	0.11	2.00	**1.00**	0.35
	1	0.21	0.99	0.35	0.16	0.12	2.00	**1.00**	0.34
	2	0.22	0.98	0.36	0.17	0.11	3.00	0.99	0.36
	3	0.19	**1.00**	0.33	0.13	0.13	**1.00**	**1.00**	0.33
	4	0.19	**1.00**	0.33	0.13	0.13	**1.00**	**1.00**	0.33
TeamZed	0	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00
	1	0.42	0.12	0.19	0.16	0.16	3.00	0.99	0.19
	2	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00
	3	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00
	4	0.00	0.00	0.00	0.17	0.17	0.00	0.00	0.00

Table 7. Task 2 ranking-based evaluation ordered by best $NDCG@100$ at 500 writings.

		1 writing			100 writings			250 writings			500 writings		
Team	Run	$P@10$	$NDCG@10$	$NDCG@100$	$P@10$	$NDCG@10$	$NDCG@100$	$P@10$	$NDCG@10$	$NDCG@100$	$P@10$	$NDCG@10$	$NDCG@100$
DeepCare	0	**1.00**	**1.00**	**0.74**	**1.00**	**1.00**	**0.88**	**1.00**	**1.00**	**0.90**	**1.00**	**1.00**	**0.91**
	1	0.90	0.93	0.65	**1.00**	**1.00**	0.79	**1.00**	**1.00**	0.83	**1.00**	**1.00**	0.85
	2	0.90	0.93	0.65	**1.00**	**1.00**	0.79	**1.00**	**1.00**	0.83	**1.00**	**1.00**	0.85
	3	0.90	0.93	0.67	**1.00**	**1.00**	0.79	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
	4	**1.00**	**1.00**	0.65	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.83	**1.00**	**1.00**	0.85
MindAILab	0	**1.00**	**1.00**	0.69	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.89	**1.00**	**1.00**	0.90
NoirByte	0	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.88	**1.00**	**1.00**	0.89
	1	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.88	**1.00**	**1.00**	0.89
	2	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.88	**1.00**	**1.00**	0.89
	3	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.88	**1.00**	**1.00**	0.89
	4	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.86	**1.00**	**1.00**	0.88	**1.00**	**1.00**	0.89
INSA-Lyon	0	0.80	0.87	0.64	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	**1.00**	**1.00**	0.89
	1	0.80	0.87	0.64	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	**1.00**	**1.00**	0.89
	2	0.90	0.81	0.55	**1.00**	**1.00**	0.80	**1.00**	**1.00**	0.83	**1.00**	**1.00**	0.85
	3	0.80	0.87	0.64	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	**1.00**	**1.00**	0.89
	4	0.80	0.87	0.63	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	**1.00**	**1.00**	0.89
LCDAA	0	0.90	0.92	0.61	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	0.90	0.94	0.88
	1	0.90	0.92	0.61	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.87	0.90	0.94	0.88
	2	0.90	0.94	0.63	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.86	0.90	0.94	0.88
	3	0.90	0.92	0.71	0.90	0.93	0.79	0.90	0.94	0.79	0.90	0.88	0.84
	4	0.90	0.93	0.73	0.90	0.93	0.78	0.90	0.92	0.79	0.90	0.86	0.81
HUTECH-NLP	0	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
	1	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
	2	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
	3	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
	4	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.84	**1.00**	**1.00**	0.85
Snugi-AI-v2	0	0.20	0.22	0.20	**1.00**	**1.00**	0.76	**1.00**	**1.00**	0.78	**1.00**	**1.00**	0.79
	1	0.20	0.22	0.20	**1.00**	**1.00**	0.76	**1.00**	**1.00**	0.78	**1.00**	**1.00**	0.79
	2	0.20	0.22	0.20	**1.00**	**1.00**	0.76	**1.00**	**1.00**	0.78	**1.00**	**1.00**	0.79
	3	0.20	0.22	0.20	**1.00**	**1.00**	0.80	**1.00**	**1.00**	0.82	**1.00**	**1.00**	0.85
	4	0.20	0.22	0.20	**1.00**	**1.00**	0.73	0.90	0.92	0.74	0.90	0.88	0.72
UNED-GELP	0	0.00	0.00	0.11	0.20	0.19	0.11	0.10	0.06	0.10	0.10	0.06	0.12
	1	0.30	0.37	0.36	0.30	0.36	0.42	0.30	0.31	0.38	0.60	0.59	0.42
	2	0.90	0.94	0.68	1.00	1.00	0.87	1.00	1.00	0.82	1.00	1.00	0.83
	3	0.90	0.94	0.69	1.00	1.00	0.88	1.00	1.00	0.85	1.00	1.00	0.84
	4	0.90	0.93	0.72	1.00	1.00	0.87	1.00	1.00	0.83	1.00	1.00	0.83
erisk-cedri	0	0.80	0.86	0.65	0.90	0.94	0.81	**1.00**	**1.00**	0.84	0.90	0.94	0.84
	1	0.70	0.72	0.60	0.90	0.93	0.77	0.90	0.94	0.77	0.90	0.94	0.76
	2	0.80	0.87	0.63	**1.00**	**1.00**	0.77	**1.00**	**1.00**	0.76	**1.00**	**1.00**	0.77
	3	0.80	0.87	0.63	**1.00**	**1.00**	0.75	**1.00**	**1.00**	0.73	**1.00**	**1.00**	0.75
	4	0.80	0.87	0.63	**1.00**	**1.00**	0.72	**1.00**	**1.00**	0.71	**1.00**	**1.00**	0.70
VANGUARD	0	1.00	1.00	0.71	0.90	0.94	0.75	0.90	0.94	0.75	0.90	0.94	0.75
	1	1.00	1.00	0.67	0.90	0.94	0.80	0.90	0.94	0.80	0.90	0.94	0.80
HUGETIME	0	0.70	0.81	0.32	0.60	0.44	0.73	0.60	0.44	0.79	0.60	0.44	0.79
Awakened	0	1.00	1.00	0.54	0.90	0.94	0.63	0.90	0.94	0.63	0.90	0.94	0.63
NYCU-NLP	0	0.60	0.62	0.46	0.70	0.79	0.46	0.80	0.76	0.47	0.80	0.76	0.51
	1	0.60	0.62	0.46	0.70	0.79	0.46	0.80	0.76	0.47	0.70	0.76	0.49
	2	0.60	0.62	0.46	0.70	0.79	0.46	0.80	0.76	0.47	0.70	0.76	0.49
	3	0.60	0.62	0.48	0.70	0.75	0.46	0.70	0.75	0.48	0.70	0.75	0.51
	4	0.60	0.62	0.48	0.70	0.75	0.46	0.70	0.75	0.48	0.70	0.75	0.51
Lotu-ixa	0	0.90	0.88	0.63	0.60	0.68	0.56	0.80	0.82	0.41	0.00	0.00	0.23
	1	0.90	0.94	0.60	0.80	0.86	0.59	0.80	0.87	0.42	0.00	0.00	0.25
	2	0.90	0.94	0.60	0.80	0.86	0.59	0.80	0.87	0.42	0.00	0.00	0.25
	3	0.90	0.94	0.60	0.80	0.86	0.59	0.80	0.87	0.42	0.00	0.00	0.25
	4	**1.00**	**1.00**	0.67	0.90	0.88	0.58	0.90	0.88	0.49	0.80	0.85	0.38
MacEarlyDetect	0	0.20	0.16	0.27	0.20	0.18	0.39	0.00	0.00	0.23	0.00	0.00	0.23
	1	0.20	0.16	0.27	0.20	0.18	0.39	0.00	0.00	0.23	0.00	0.00	0.23
	2	0.20	0.16	0.27	0.20	0.18	0.39	0.00	0.00	0.23	0.00	0.00	0.23
	3	0.50	0.58	0.37	0.90	0.94	0.55	0.00	0.00	0.25	0.00	0.00	0.25
	4	0.50	0.58	0.37	0.90	0.94	0.55	0.00	0.00	0.25	0.00	0.00	0.25
UET-PsyWar	0	0.60	0.71	0.32	0.50	0.64	0.30	0.00	0.00	0.23	0.00	0.00	0.23
	1	0.60	0.61	0.30	0.50	0.64	0.30	0.00	0.00	0.23	0.00	0.00	0.23
	2	0.60	0.60	0.30	0.50	0.64	0.30	0.00	0.00	0.23	0.00	0.00	0.23
	3	0.60	0.60	0.30	0.50	0.64	0.30	0.00	0.00	0.23	0.00	0.00	0.23
	4	0.60	0.61	0.30	0.50	0.64	0.30	0.00	0.00	0.23	0.00	0.00	0.23
TeamZed	0	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20
	1	0.50	0.57	0.30	0.50	0.57	0.30	0.50	0.57	0.30	0.50	0.57	0.30
	2	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20
	3	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20
	4	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20	0.20	0.22	0.20

3 Task 3: ADHD Symptom Sentence Ranking

In 2026, we introduced a new sentence-level ranking task focused on the retrieval of evidence related to ADHD symptoms. This task extends the symptom-oriented retrieval setting explored in previous eRisk editions. In this case, we target the ADHD disorder, and the 18 symptoms defined in the Adult ADHD Self-Report Scale (ASRS-v1.1). Participants were asked to rank candidate sentences according to their relevance to each ADHD symptom. A sentence was considered relevant when it conveyed information about the user's state with respect to the target symptom, independently of polarity or stance. Therefore, both sentences expressing the presence of a symptom and sentences providing evidence about its absence could be relevant, as long as they were informative for the corresponding symptom.

The task was designed to encourage systems to identify clinically meaningful evidence at sentence granularity, rather than relying only on surface keyword matching. This is especially important in ADHD-related language, where symptoms such as inattention, impulsivity, disorganization, restlessness, or difficulty sustaining effort may be expressed indirectly, through everyday experiences and self-descriptions. By formulating the task as a ranking problem, participants were required to prioritize the most informative candidate sentences for each symptom.

As this was the first edition of the ADHD symptom ranking task, no annotated training data was provided. Participants received the test inputs only, following the formatting conventions of recent eRisk ranking tasks. For each of the 18 ASRS-v1.1 symptoms, teams submitted a ranking of up to 1000 candidate sentences, ordered by decreasing estimated relevance. Relevance assessments were then produced using top-k pooling and expert annotation, and systems were evaluated using standard information retrieval metrics such as MAP, nDCG, and P@10.

3.1 Task 3: Dataset

The dataset released for this task consisted of sentence-tagged documents derived from publicly available social media writings collected to contain ADHD-related expressions. The corpus was provided in a format compatible with previous eRisk ranking tasks, where each sentence receives a unique identifier that participants use in their submitted rankings. Table 8 presents the main statistics of the corpus.

Table 8. Corpus statistics for Task 3: ADHD Symptom Sentence Ranking.

Number of sentences	4 170 875
Average number of words per sentence	13.05
Number of ADHD symptoms	18

3.2 Task 3: Assessment Process

Participants were given the corpus of candidate sentences and the description of the 18 symptoms from the ASRS-v1.1 questionnaire. They were free to decide the best strategy to represent each ADHD symptom as a query or retrieval target. Each team could submit up to 5 variants, or runs. Each run included 18 TREC-style formatted rankings of sentences, one ranking per ADHD symptom, as illustrated in Figure 2. For each symptom, participants submitted up to 1000 results sorted by estimated relevance.

```
1   Q0  sent_000251_0    0001    10.0    myGroupNameMyMethodName
1   Q0  sent_085820_3    0002     9.4    myGroupNameMyMethodName
...
18  Q0  sent_015320_2    0998     1.2    myGroupNameMyMethodName
18  Q0  sent_022313_9    1000     0.9    myGroupNameMyMethodName
```

Figure 2. Example of a participant's run for Task 3.

The relevance assessment process followed the same methodology used in recent eRisk sentence-ranking tasks. For each ADHD symptom, a pool of candidate sentences was created from the highest-ranked results submitted by the participating systems. These pooled sentences were then manually assessed by expert annotators. A sentence was judged as relevant when it provided evidence about the user's state with respect to the target ADHD symptom, regardless of whether the sentence indicated the presence, absence, uncertainty, or contextual manifestation of that symptom.

All sentences included in the assessment pool were independently labelled by three expert assessors. From these annotations, we derived two versions of the relevance judgments. In the *majority-vote* version, a sentence was considered relevant if at least two of the three assessors judged it as relevant. In the *unanimity* version, a sentence was considered relevant only if all three assessors agreed on its relevance. Therefore, the unanimity version constitutes a more restrictive evaluation setting, focusing only on sentences with full agreement among assessors.

The resulting relevance judgments were used to evaluate the submitted rankings with standard information retrieval metrics. In particular, we report Mean Average Precision (MAP), normalized Discounted Cumulative Gain (nDCG), including nDCG@100, and Precision at 10 (P@10). These measures capture complementary aspects of ranking quality: MAP rewards systems that retrieve relevant sentences throughout the ranking, nDCG emphasizes the ordering of relevant evidence near the top of the ranking, and P@10 reflects the quality of the highest-ranked sentences that would be most immediately inspected by a user or expert. We received 45 runs from 12 participating teams, as shown in Table 9.

Table 9. Task 3: participating teams and number of submitted runs for the ADHD Symptom Sentence Ranking task.

Team	#Runs
AI-MH-Z	4
Awakened	2
DS-GT	5
INSA-Lyon	5
NaiveNeuron	5
NeuroRankUB	4
PMH-PsychAI-Lab	1
Silvermask	1
SymptomAI	4
UET-PsyWar	5
VANGUARD	4
erisk-cedri	5
Total	45

3.3 Task 3: Results

The submitted runs were evaluated under the two relevance-judgment settings derived from the expert annotations. Table 10 reports the results obtained with the majority-vote relevance judgments, where a sentence is considered relevant when at least two of the three assessors marked it as relevant. Table 11 reports the results obtained with the unanimity relevance judgments, where only sentences marked as relevant by all three assessors are considered relevant. The unanimity setting is therefore more restrictive and provides a complementary view of system performance on sentences with stronger assessor agreement.

Table 10. Task 3 results using majority-vote relevance judgments. A sentence is considered relevant if at least two of the three expert assessors marked it as relevant.

Team	Run	AP	R-PREC	P@10	NDCG
NeuroRankUB	final ADHD ranking	0.248	0.328	0.528	0.547
NeuroRankUB	gte rrf final	0.203	0.264	0.450	0.516
erisk-cedri	EnsembleV3	0.190	0.262	0.400	0.461
erisk-cedri	EnsembleV2	0.186	0.259	0.406	0.459
NeuroRankUB	ADHD submission	0.171	0.242	0.417	0.475
erisk-cedri	EnsembleV1	0.167	0.239	0.372	0.446
INSA-Lyon	Ensemble	0.159	0.206	0.317	0.406
INSA-Lyon	LLM cascade	0.158	0.180	0.344	0.402
INSA-Lyon	HiPerT full	0.137	0.169	0.294	0.388
INSA-Lyon	BiEnc baseline	0.117	0.160	0.311	0.321
erisk-cedri	VecSearch	0.105	0.162	0.278	0.350
erisk-cedri	NLI zero-shot	0.097	0.162	0.267	0.379
NeuroRankUB	final ranking	0.092	0.185	0.283	0.312
DS-GT	asrs run5	0.085	0.115	0.300	0.207
AI-MH-Z	ST equal weights	0.076	0.110	0.239	0.317
AI-MH-Z	ST option weights	0.075	0.110	0.228	0.316
AI-MH-Z	ST variable weights	0.074	0.109	0.228	0.313
UET-PsyWar	run5 RRF robust	0.072	0.134	0.272	0.270
UET-PsyWar	run4 hybrid key bonus	0.069	0.121	0.222	0.267
UET-PsyWar	run3 balanced hybrid	0.068	0.121	0.222	0.264
SymptomAI	run4	0.067	0.096	0.300	0.215
INSA-Lyon	DepTransfer	0.054	0.087	0.100	0.216
SymptomAI	run1	0.052	0.098	0.194	0.267
UET-PsyWar	run2 CE only	0.050	0.107	0.200	0.236
Silvermask	v2 semantic	0.045	0.090	0.167	0.241
NaiveNeuron	ENSEMBLE	0.044	0.094	0.144	0.220
DS-GT	asrs run4	0.043	0.078	0.194	0.146
DS-GT	asrs run2	0.038	0.075	0.189	0.134
NaiveNeuron	RERANK MultiQA	0.037	0.090	0.139	0.204
SymptomAI	run3	0.037	0.056	0.189	0.155
AI-MH-Z	ST by description	0.035	0.062	0.094	0.150
NaiveNeuron	Rerank BioMed	0.034	0.081	0.150	0.152
NaiveNeuron	HYBRID RRF	0.031	0.073	0.128	0.127
PMH-PsychAI-Lab	run1	0.029	0.075	0.117	0.189
UET-PsyWar	run1 sim only	0.029	0.072	0.106	0.172
VANGUARD	run3	0.028	0.065	0.156	0.113
DS-GT	asrs run1	0.026	0.044	0.094	0.109
DS-GT	asrs run3	0.023	0.058	0.122	0.112
VANGUARD	run5	0.023	0.049	0.111	0.106
VANGUARD	run1	0.014	0.033	0.056	0.089
VANGUARD	run4	0.014	0.030	0.083	0.052
SymptomAI	run2	0.012	0.041	0.056	0.094
Awakened	run1	0.006	0.015	0.022	0.057
Awakened	run2	0.006	0.019	0.028	0.058
NaiveNeuron	S-Biomed-Roberta	0.006	0.011	0.022	0.095

Table 11. Task 3 results using unanimity relevance judgments. A sentence is considered relevant only if all three expert assessors marked it as relevant.

Team	Run	AP	R-PREC	P@10	NDCG
NeuroRankUB	final ADHD ranking	0.207	0.274	0.367	0.488
NeuroRankUB	gte rrf final	0.184	0.233	0.306	0.466
NeuroRankUB	ADHD submission	0.166	0.232	0.311	0.439
erisk-cedri	EnsembleV3	0.138	0.193	0.261	0.387
INSA-Lyon	LLM cascade	0.134	0.167	0.272	0.369
erisk-cedri	EnsembleV2	0.134	0.191	0.278	0.384
INSA-Lyon	Ensemble	0.129	0.160	0.244	0.363
INSA-Lyon	BiEnc baseline	0.122	0.154	0.256	0.313
erisk-cedri	EnsembleV1	0.122	0.190	0.233	0.372
INSA-Lyon	HiPerT full	0.107	0.131	0.172	0.336
erisk-cedri	VecSearch	0.090	0.148	0.200	0.304
DS-GT	asrs run5	0.080	0.121	0.217	0.188
UET-PsyWar	run4 hybrid key bonus	0.079	0.122	0.156	0.276
AI-MH-Z	ST equal weights	0.078	0.108	0.167	0.291
AI-MH-Z	ST option weights	0.078	0.105	0.161	0.289
AI-MH-Z	ST variable weights	0.076	0.102	0.167	0.288
NeuroRankUB	final ranking	0.073	0.117	0.183	0.268
UET-PsyWar	run5 RRF robust	0.072	0.113	0.200	0.262
UET-PsyWar	run3 balanced hybrid	0.071	0.128	0.167	0.268
SymptomAI	run4	0.070	0.093	0.206	0.192
UET-PsyWar	run2 CE only	0.058	0.099	0.161	0.242
erisk-cedri	NLI zero-shot	0.057	0.093	0.117	0.297
INSA-Lyon	DepTransfer	0.054	0.065	0.078	0.202
DS-GT	asrs run4	0.052	0.078	0.133	0.142
NaiveNeuron	ENSEMBLE	0.049	0.093	0.117	0.208
DS-GT	asrs run2	0.048	0.065	0.144	0.135
SymptomAI	run1	0.048	0.072	0.133	0.234
NaiveNeuron	Rerank BioMed	0.043	0.093	0.122	0.159
NaiveNeuron	RERANK MultiQA	0.042	0.082	0.111	0.191
Silvermask	v2 semantic	0.041	0.070	0.094	0.199
DS-GT	asrs run1	0.040	0.057	0.083	0.116
NaiveNeuron	HYBRID RRF	0.037	0.076	0.089	0.133
VANGUARD	run3	0.034	0.071	0.117	0.115
AI-MH-Z	ST by description	0.031	0.051	0.072	0.129
PMH-PsychAI-Lab	run1	0.030	0.066	0.083	0.180
SymptomAI	run3	0.030	0.040	0.128	0.126
VANGUARD	run5	0.029	0.051	0.078	0.110
DS-GT	asrs run3	0.026	0.053	0.083	0.100
UET-PsyWar	run1 sim only	0.020	0.044	0.050	0.149
VANGUARD	run1	0.018	0.031	0.039	0.090
VANGUARD	run4	0.018	0.027	0.044	0.056
Awakened	run1	0.009	0.017	0.017	0.054
NaiveNeuron	S-Biomed-Roberta	0.006	0.010	0.017	0.090
Awakened	run2	0.005	0.011	0.017	0.043
SymptomAI	run2	0.005	0.015	0.028	0.066

References

1. David E. Losada and Fabio Crestani. A test collection for research on depression and language use. In *Proceedings Conference and Labs of the Evaluation Forum CLEF 2016*, Evora, Portugal, 2016.
2. David E Losada, Fabio Crestani, and Javier Parapar. Overview of erisk 2019 early risk prediction on the internet. In *Experimental IR Meets Multilinguality, Multimodality, and Interaction: 10th International Conference of the CLEF Association, CLEF 2019, Lugano, Switzerland, September 9–12, 2019, Proceedings 10*, pages 340–357. Springer, 2019.
3. Farig Sadeque, Dongfang Xu, and Steven Bethard. Measuring the latency of depression detection in social media. In *WSDM*, pages 495–503. ACM, 2018.
4. Marcel Trotzek, Sven Koitka, and Christoph M. Friedrich. Utilizing neural networks and linguistic metadata for early detection of depression indications in text sequences. *CoRR*, abs/1804.07000, 2018.